

16831 Statistical Techniques, Fall 2011

Homework 5: Online Learning

Natasha Kholgade and Marynel Vázquez

December 7th, 2011

Classifiers

For this homework, we implemented Gaussian Process Regression (GPR), Boosting, and online linear SVMs. We used these algorithms in a binary fashion to classify 3D points into five different classes: Vegetation, Wire, Pole, Ground and Facade. Each classifier was trained 5 times in a one-vs-all scenario, such that the final classification for a testing data point was given by the class with highest score.

GPR and boosting were not implemented in an online fashion (though boosting does lend itself to an online implementation, we were limited in time). Our version of Gaussian Process Regression uses the exponential to the negative squared distance between features (radial basis function) as the covariance function. Meanwhile, boosting performs feature selection, and uses the exponentiated gradient as suggested by Alex Grubb.

Our implementation of a Support Vector Machine is based on the subgradient of the hinge loss, plus a regularization term. It operates fast, though requires a significant number of samples to satisfactorily classify new data in comparison to the other algorithms being considered.

Performance Summary

For this homework, we performed two-fold cross-validation on data from one file, and then tested with data from the other file. We swapped the two files, and re-did cross-validation and training. We provide the best parameters, confusion matrices, per-class percentage performance, and net classification rate for the three classes.

For brevity, we refer to the point cloud data file *oakland_part3_am_rf.node_features* as *File_am*, and to the file *oakland_part3_an_rf.node_features* as *File_an* in the following sections.

Gaussian Process Regression

Train with first file, test with the second file:

Best parameters: radial basis function parameter $\sigma = .4$, regularization parameter $\lambda = .4642$.

Net classification rate: .8637

Per-class classification rate:

[0.8208 0.8870 0.7192 0.9796 0.6670]

Confusion matrix:

$$\begin{bmatrix} 0.8208 & 0.0560 & 0.0407 & 0.0015 & 0.0809 \\ 0.0758 & 0.8870 & 0.0084 & 0.0004 & 0.0285 \\ 0.1971 & 0.0092 & 0.7192 & 0.0046 & 0.0700 \\ 0.0018 & 0.0179 & 0.0003 & 0.9796 & 0.0005 \\ 0.1316 & 0.1576 & 0.0421 & 0.0017 & 0.6670 \end{bmatrix}$$

Train with first file, test with the second file:

Best parameters: radial basis function parameter $\sigma = .4$, regularization parameter $\lambda = .4642$.

Net classification rate: .947

Per-class classification rate:

[0.8097 0.5697 0.9279 0.9899 0.8305]

Confusion matrix:

[0.8097 0.0471 0.0680 0.0010 0.0743
0.0685 0.5697 0.0183 0.0697 0.2738
0.0378 0.0014 0.9279 0 0.0329
0.0020 0.0001 0.0065 0.9899 0.0014
0.0611 0.0391 0.0686 0.0007 0.8305]

Linear SVMs

AdaBoost

Train with first file, test with the second file:

Best parameters: radial basis function parameter $\sigma = .4$, regularization parameter $\lambda = .4642$.

Net classification rate: .8637

Per-class classification rate:

[0.8208 0.8870 0.7192 0.9796 0.6670]

Confusion matrix:

[0.8208 0.0560 0.0407 0.0015 0.0809
0.0758 0.8870 0.0084 0.0004 0.0285
0.1971 0.0092 0.7192 0.0046 0.0700
0.0018 0.0179 0.0003 0.9796 0.0005
0.1316 0.1576 0.0421 0.0017 0.6670]

Train with first file, test with the second file:

Best parameters: radial basis function parameter $\sigma = .4$, regularization parameter $\lambda = .4642$.

Net classification rate: .947

Per-class classification rate:

[0.8097 0.5697 0.9279 0.9899 0.8305]

Confusion matrix:

[0.8097 0.0471 0.0680 0.0010 0.0743
0.0685 0.5697 0.0183 0.0697 0.2738
0.0378 0.0014 0.9279 0 0.0329
0.0020 0.0001 0.0065 0.9899 0.0014
0.0611 0.0391 0.0686 0.0007 0.8305]

We find that points belonging to the 'Pole' and 'Ground' classes do not perform well because there is not enough data to train them.

GPR and SVM were relatively easy to implement. AdaBoost required a bit of delving into some math, but we made it through.

Supper Vector Machine

Train with first file, test with the second file:

Train with first file, test with the second file:

1. How well did it perform for online learning? Does it perform well on the held-out data?
2. Are there any classes that did not get classified well? Why do you think that is?
3. How easy was the learner to implement?
4. How long does the learner take (in terms of data points, dimensions, classes, etc...) for training and prediction?
5. Show images/movies of the classified data. Note that MATLAB is not very good at displaying thousands of 3D points; use VRML or python.
6. How did you choose (hyper)parameters (priors, kernel width, noise variance, prior variance, learning rate, etc...)?
7. How robust is this algorithm to noise? Take the current feature set and:
 - Add a large number of random features
 - Add a large number of features that are noise corrupted versions of the features already in the data-set.

You should also compare the learners' performance to each other. Did kernels help on this data set? Which one would you use on your robot? What would future work include?